ML Project Documentation

Linear Regression

1. ***General Information on dataset:***
    a. *The name of dataset used: California Housing Data (1990).*
    b. *The number of samples used in training and testing:*
        i. Training : `14087`
        ii. Testing : `3522`
2. ***Implementation details:***
    a. *At feature extraction phase, how many features were extracted, their names, the dimension of resulted features.*
        i. **Number of features extracted:** After preprocessing and feature engineering, the dataset contains additional features.
        ii. **Names of resulted features:** `rooms_per_household`, `bedrooms_per_room` and `population_per_household`.
        iii. **Dimension of resulted features:** `(17609, 16)`
    b. *Is cross-validation is used in any of implemented models? Yes.*
        i. *The number of fold:* `3-fold cross-validation`.
        ii. *Ratio of training/validation.*
            1. Training set: 2/3 of the data
            2. Validation set: 1/3 of the data
3. ***Results details:***

    1.*Loss curve*

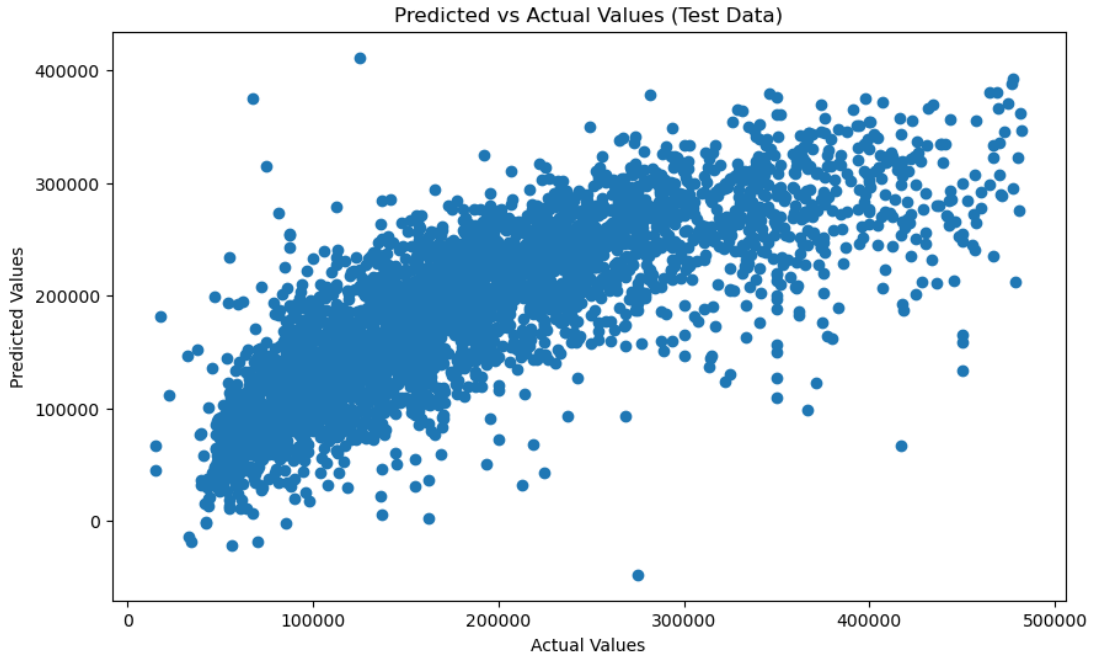

2. *Accuracy*

```
Average R squared score :  0.6231679206055769
```

```
Evaluation score on 3 cross-validation sets :  [0.6145668  0.62858995
0.62634702]
```

KNN

1. ***General Information on dataset:***
    a. *the name of dataset used:*
    b. *the number of samples used in training and testing.*
        i. Training: `14087`
        ii. Testing: `3522`
2. ***Implementation details:***

- *At feature extraction phase, how many features were extracted, their names, the dimension of resulted features.*
    - **Number of features extracted:** After preprocessing and feature engineering, the dataset contains additional features.

- **Names of resulted features:** `rooms_per_household`, `bedrooms_per_room` and `population_per_household`.
- **Dimension of resulted features:** `(17609, 16)`

- *Is cross-validation is used in any of implemented models? Yes.*
  1. *The number of fold:* `3-fold cross-validation`.
  2. *Ratio of training/validation.*
     a. Training set: 2/3 of the data
     b. Validation set: 1/3 of the data
- *Hyperparameters used in your model, as initial learning rate, optimizer, regularization, batch size, no. of epochs, etc…*
  1. **K-Nearest Neighbors (KNN):**
     - `n_neighbors`: Number of neighbors to consider `18 in your example`.
  2. **Gradient Boosting Regressor:**
     - `n_estimators`: Number of boosting stages `100`.
     - `learning_rate`: The step size shrinkage used to prevent overfitting `0.1`.
     - `max_depth`: Maximum depth of the individual trees `8 in your example`.
  3. **XGBoost:**
     - `max_depth`: Maximum depth of a tree `8`.
     - `n_estimators`: Number of boosting rounds `100`.
     - `objective`: The learning task and corresponding objective, set to `reg:squarederror` for regression.
     - `random_state`: Seed for reproducibility `42`.

1. ***Results details:***

   KNN Loss curve

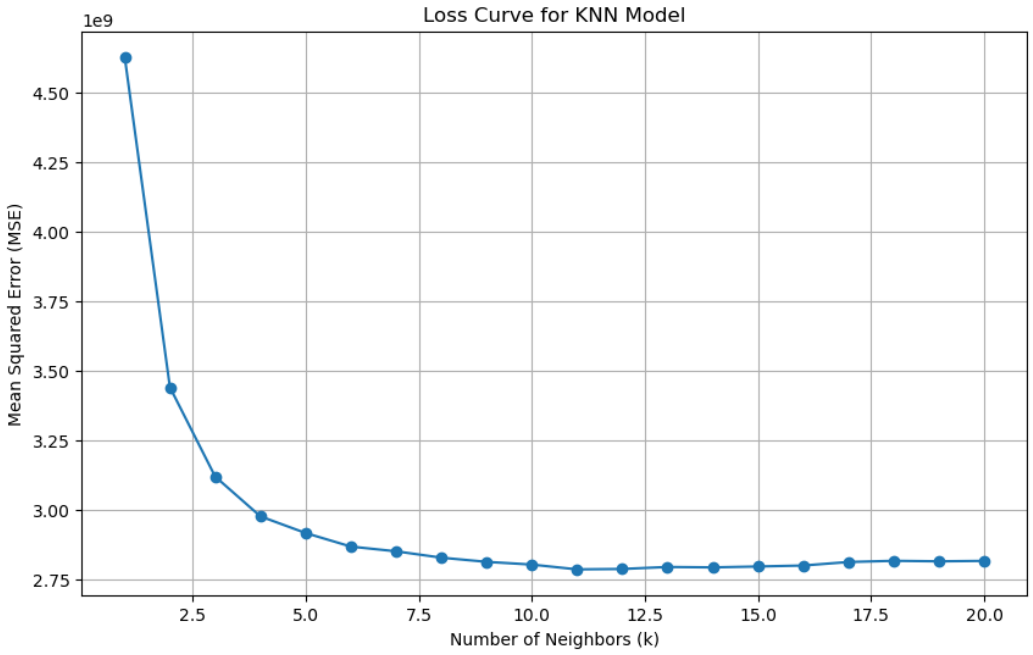

KNN Accuracy

```
Evaluation score on 3 cross-validation sets :  [0.67168832 0.68923142
0.68611015]
```

```
Average R squared score :  0.6823432959082879
```

Loss Curve For Gradient Boosting Regressor.

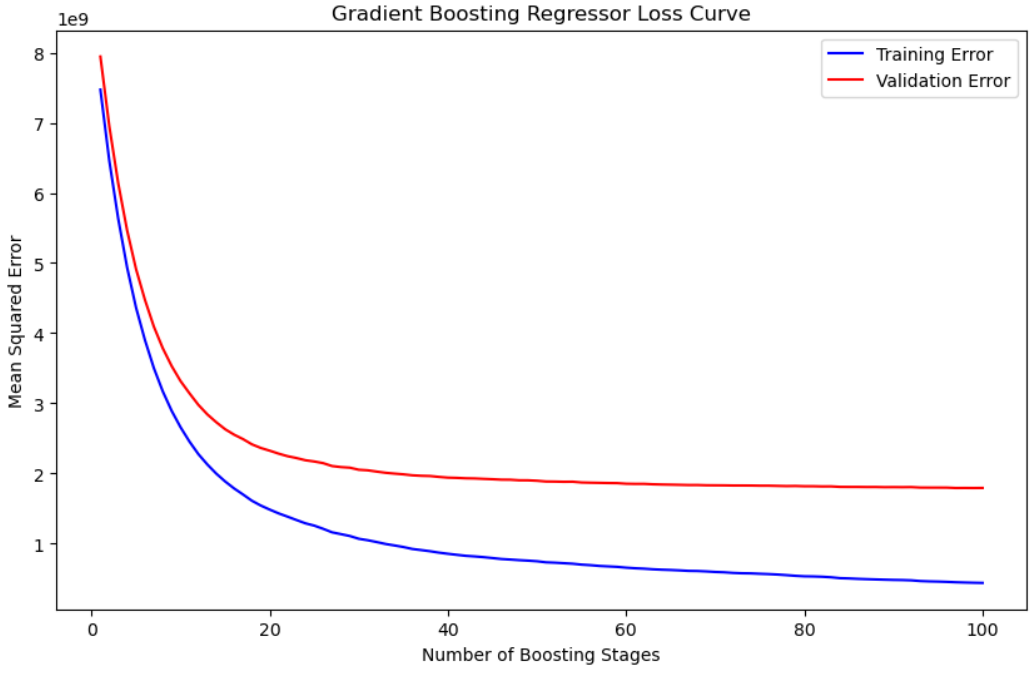

Gradient Boosting Regressor Accuracy.

```
Evaluation score on 3 cross-validation sets :  [0.78743338 0.80353852
0.80631629]
```

```
Average R squared score :  0.7990960650730976
```

Loss Curve For **eXtreme Gradient Boosting Regressor (XGBoost)**

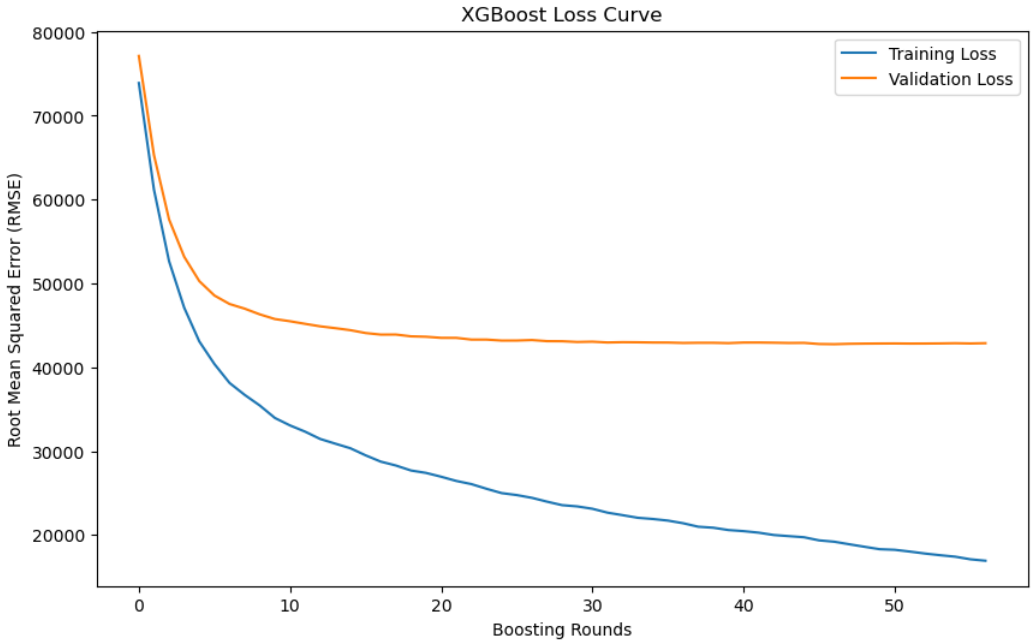

eXtreme Gradient Boosting Regressor (XGBoost) ***Accuracy***

```
Average R squared score :  0.7906049245321869
```

```
Evaluation score on 3 cross-validation sets :  [0.78188369 0.79387968
0.79605141]
```

Logistic Regression

1. ***General Information on dataset:***
    a. *the name of dataset used: Food-101*
    b. *number of classes and their labels:*

i. Five classes.

ii. ['cheesecake', 'cup_cakes', 'donuts', 'hamburger', 'pizza']

c. *the total number of samples in dataset and the size of each (in case of images):*

i. number of samples: 3500 samples

ii. Images size: (224 * 224)

d. *the number of samples used in training, validation and testing:*

i. training = `2800`

ii. validation = `0.2` of training set (560 samples).

iii. testing = `700`

2. ***Implementation details:***

- *At feature extraction phase, how many features were extracted, their names, the dimension of resulted features.*
  - *how many features were extracted:* 128 features were extracted for each image using the SIFT.
  - *the dimension of resulted features:* `(2800, 224, 224, 3)`
- *Is cross-validation is used in any of implemented models? No.*

1. ***Results details:***

Accuracy

`Accuracy: 0.22714285714285715`

Confusion matrix

```
Confusion Matrix:
[[38 31 22 26 30]
 [23 24 29 33 24]
 [30 32 29 34 27]
 [31 17 20 40 14]
 [29 25 27 37 28]]
```

ROC curve

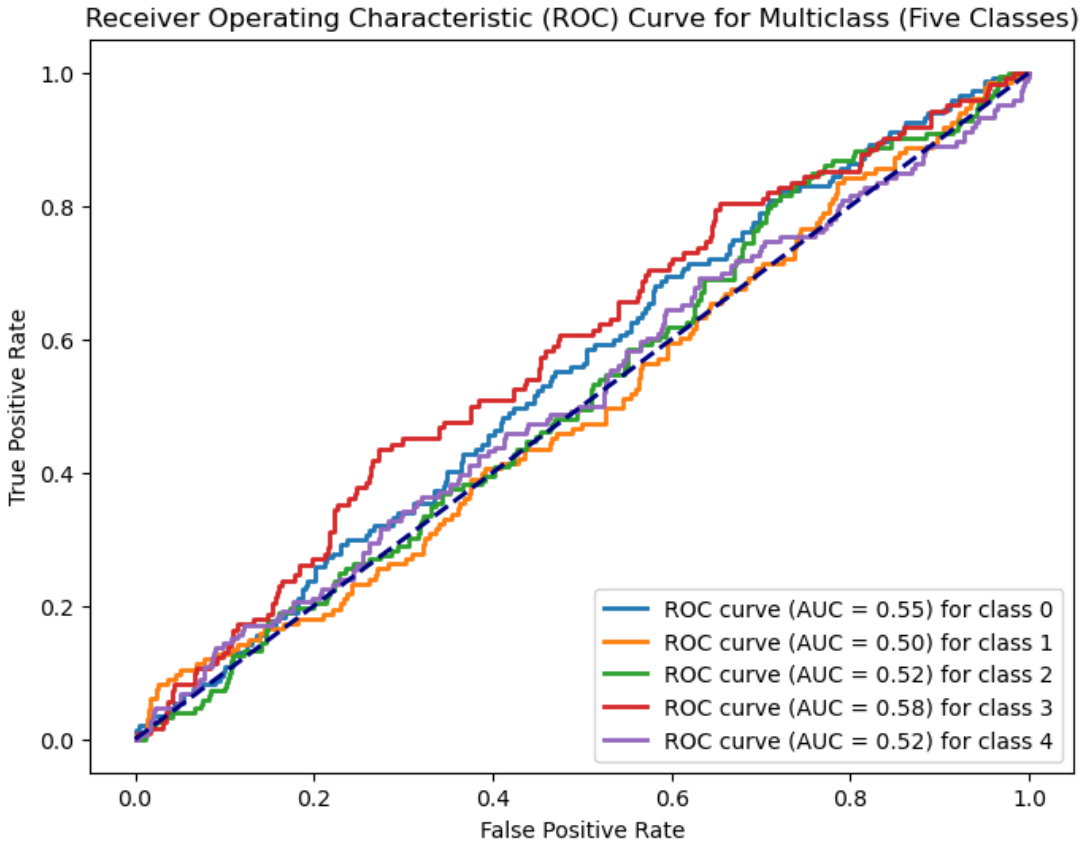
Receiver Operating Characteristic (ROC) Curve for Multiclass (Five Classes)
True Positive Rate
False Positive Rate
0.0
0.2
0.4
0.6
0.8
1.0
ROC curve (AUC = 0.55) for class 0
ROC curve (AUC = 0.50) for class 1
ROC curve (AUC = 0.52) for class 2
ROC curve (AUC = 0.58) for class 3
ROC curve (AUC = 0.52) for class 4